



---

**Spatial Reasoning in Multimodal LLMs  
via CoT Distillation and Monte Carlo Tree Search  
for Dutch Facade-Element Detection: An Exploratory Study**

---

Thesis submitted to Utrecht University for the degree of  
MSc Artificial Intelligence, July 2025.

**Riccardo Campanella**

8175721

**Examiners:** Dr. Metehan Doyran, Dr. Itir Önal Ertuğrul  
**External Supervisor:** Raphaël Gueulet (TNO)

# Presentation Outline

- ① Motivation & Problem Statement
- ② Research Questions & Contributions
- ③ Background
- ④ Data Collection & Methodology
- ⑤ Exploratory Results Analysis
- ⑥ RQ Discussion & Limitations
- ⑦ Conclusions & Future Work

# Motivation: Building Energy Renovation

## The Challenge

- Dutch buildings need energy retrofits for sustainability goals
- Manual facade assessment is time-consuming and expensive
- Current Computer Vision models lack domain-specific knowledge

## The Opportunity

- Multimodal LLMs offer contextual reasoning capabilities
- Can process natural language descriptions of architectural features
- Potential for zero-shot transfer to specialized domains

# Problem Statement

## MQA Target Features:

- 1 Weep holes
- 2 Crawling space
- 3 Chimneys
- 4 Pitched roof
- 5 Facade ventilation
- 6 Roof ventilation
- 7 Window ventilation
- 8 Dormers
- 9 Roof windows
- 10 Attics living spaces
- 11 Vegetation growth
- 12 Balconies
- 13 Photovoltaic panels
- 14 Window count



## Reasoning Complexity:

Visual Recognition → Geometric Inference →  
Semantic Understanding → Context Analysis

# Research Questions

## Main Research Question

Are SoTA\* Multimodal LLMs beneficial to identify applicable housing renovation concepts on Dutch building facades?

## RQ1: Model Comparison

How does Chain-of-Thought reasoning (Qwen) compare to 3D scene graph methods (SpatialRGPT) in zero-shot prediction?

- Performance vs. SoTA models (GPT-4o)
- Impact of bounding box guidance

## RQ2: Enhancement Methods

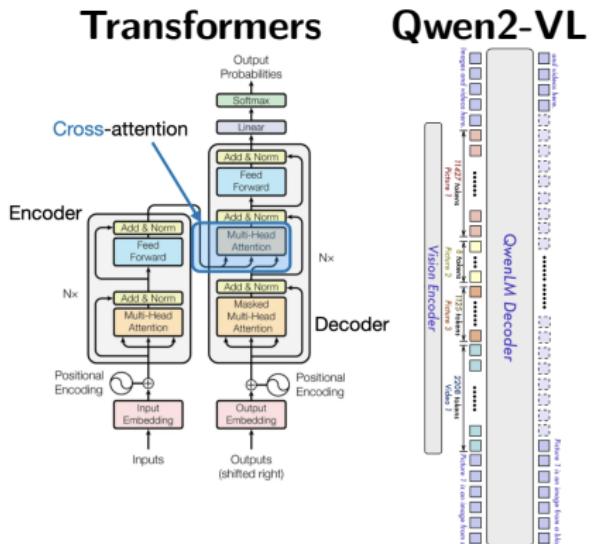
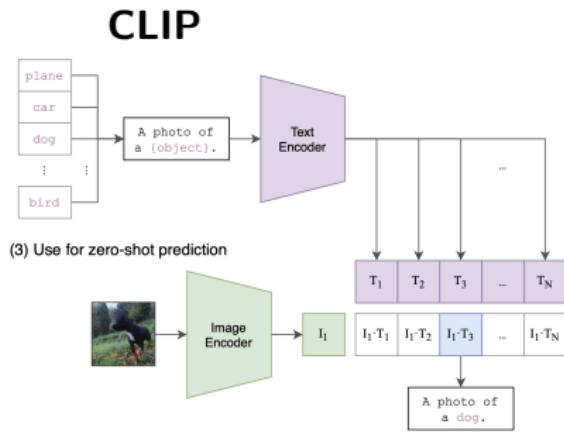
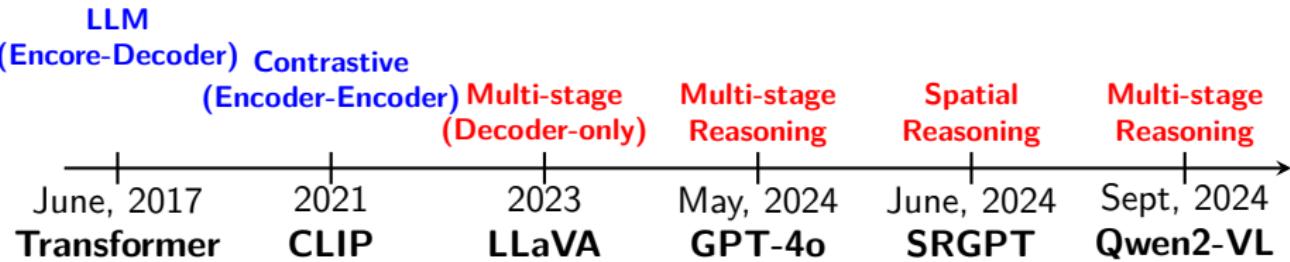
How can CoT reasoning MLLMs be enhanced for spatial recognition?

- Effect of 3D scene graph augmentation
- LoRA fine-tuning capabilities

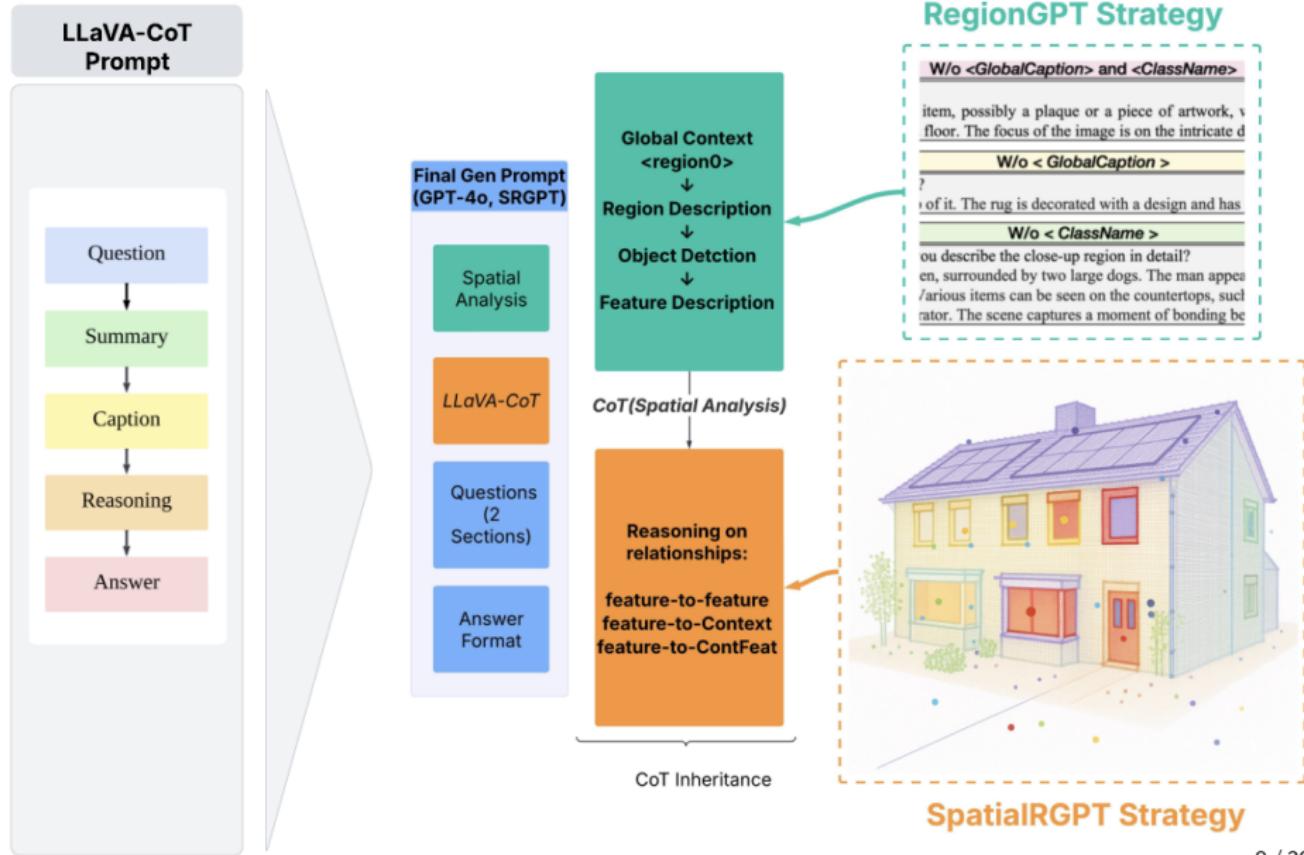
# Key Contributions

- ① **Comprehensive MLLM Evaluation Framework**
  - **First** systematic evaluation on real Dutch facade data
  - **Baseline performance** insights for architectural features
- ② **DuTCh SpaCE: Spatial Reasoning Enhancement**
  - **Mitigation:** Novel dual-teacher distillation framework to mitigate Hallucinations
  - **Compensation:** Reasoning compensates for limited visual grounding
- ③ **Reasoning vs. Grounding Trade-off Analysis**
  - **Knowledge Transfer:** Domain expertise vs. model scale
  - **Fast Scaling:** Accessible path to spatial reasoning enhancement
- ④ **Practical Domain Adaptation Insights**
  - LoRA + Knowledge Distillation + Test-time Search
  - **Few-shot:** Specialization framework for low-data domains

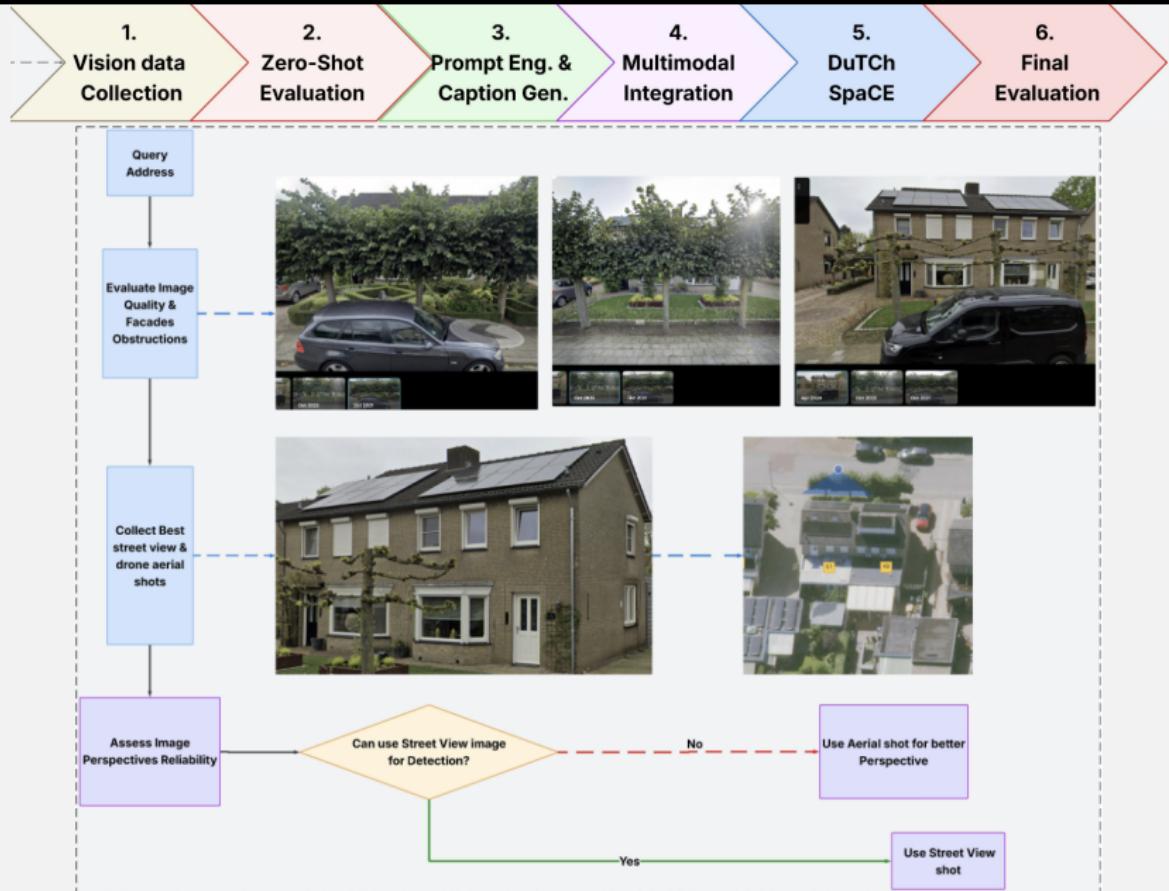
# Background: Multimodal Large Language Models



# Background: Grounding Chain-of-Thought Reasoning

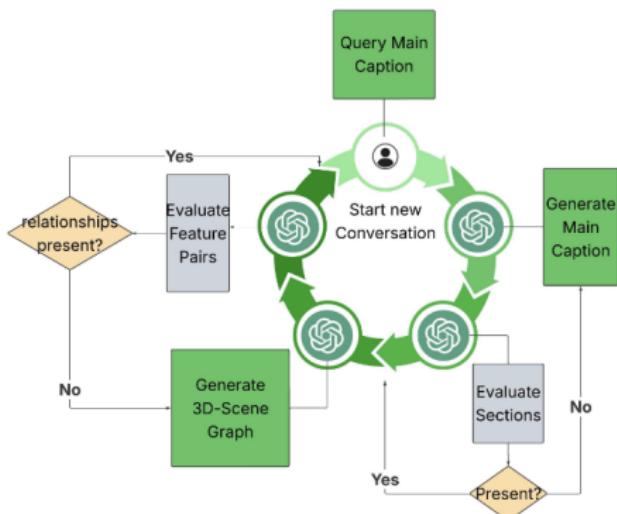


# Data Collection: Few-shot Dutch Building Facade

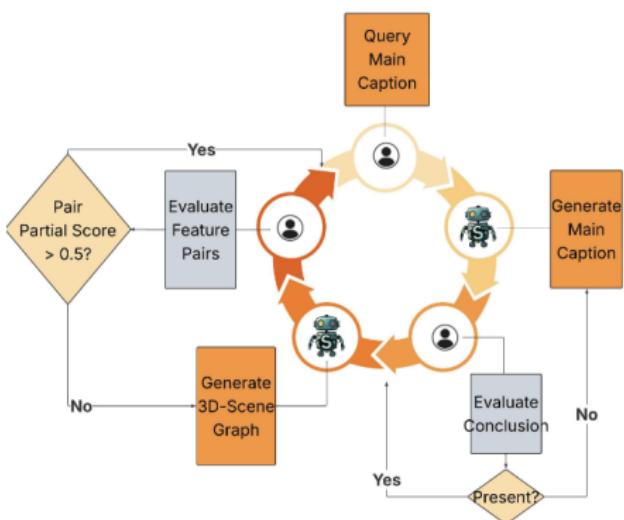


# Multimodal integration: Captions and Scenes

Teacher 1: GPT-4o



Teacher 2: SpatialIRGPT-bbox



Train Set (32)

Label	Distribution
No	44.7%
Yes	32.5%
Unknown	22.9%

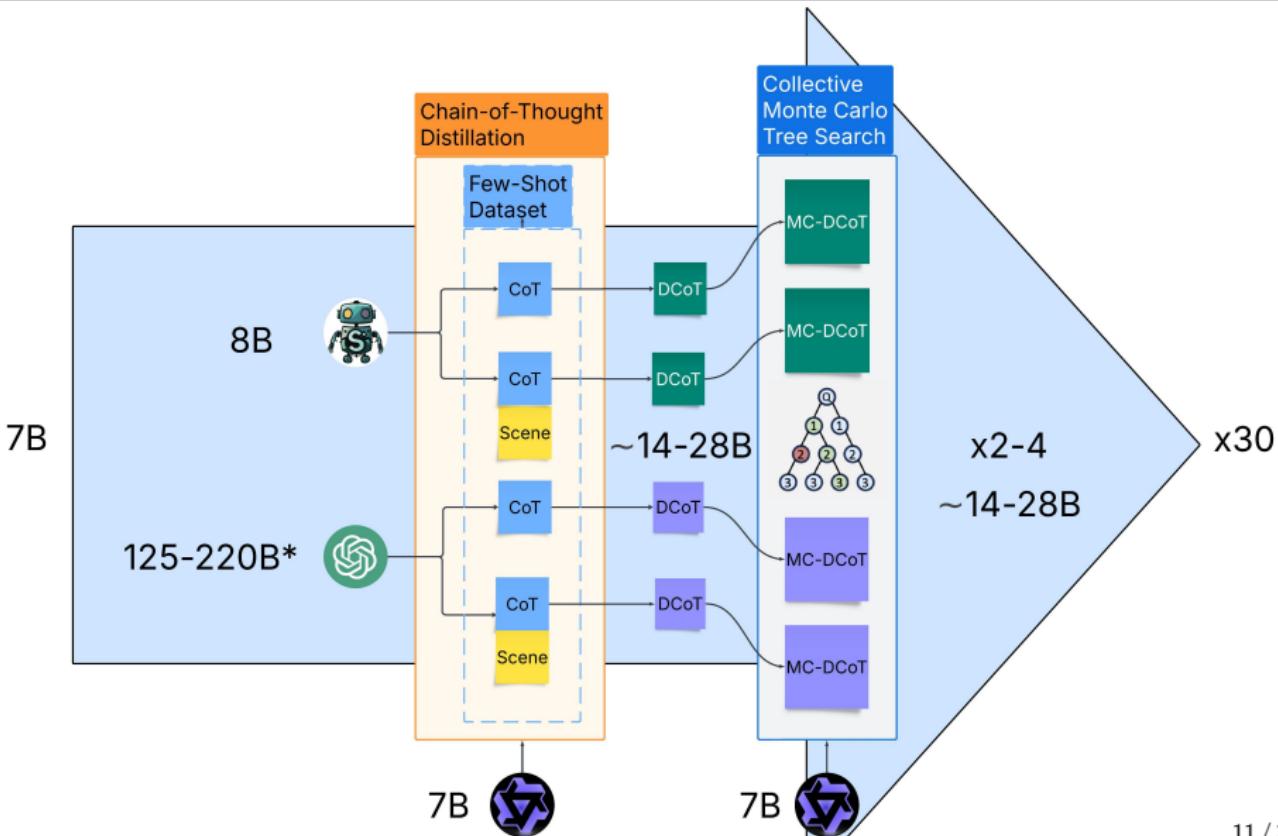
Val Set (4)

Label	Distribution
No	40.4%
Yes	46.2%
Unknown	13.5%

Test Set (9)

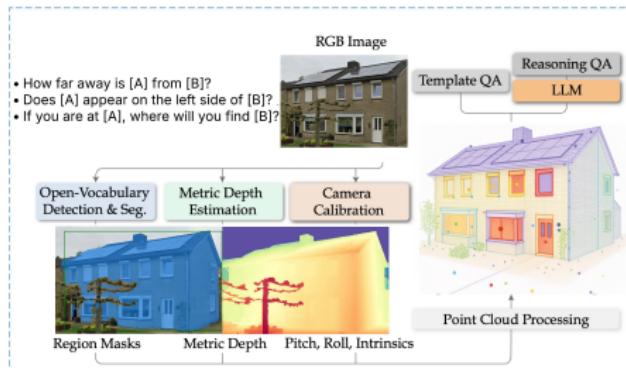
Label	Distribution
No	44.4%
Yes	39.3%
Unknown	16.2%

# Dual-Teacher Chain-of-Thought Framework for Spatial Facade Reasoning (DuTCh Space)



# Methodology: Models Under Study

- ① GPT-4o
- ② SpatialRGPT-8B (Base/Bbox)
- ③ Qwen2-7B-VL-Instruct



## Categories | Datasets

Captioning	Image Paragraph Captioning [77], MSR-VTT [78], TextCaps [79]
Reasoning	CLEVR [80], NLVR [81], VisualMRC [82]
Translation	Multi30k [83]
VQA	ActivityNet-QA [84], DocVQA [85], GQA [86], iVQA [87], MSRVTT-QA [78], MSVD-QA [78], OCR-VQA [88], ST-VQA [89], ViQuAE [90], VQAv2 [91], Visual Dialog [92]

## Categories | Datasets

Classification	V3Det [93], COCO [40], LVIS [94]
Caption	V3Det [93] VG [95], RefCOCO [96]
Relationship	VG [95]
REC	RefCOCO [96]

Table 11: The region-level SFT blend [19] we used.

User: Can you tell me if this house has photovoltaic panels?

SpatialRGPT: Yes, Ground Truth (GT)=Yes

Two photovoltaic arrays are visible on the pitched roof. The surface geometry and inclination match standard PV module dimensions (approx. 1m x 1.7m each), and they are aligned symmetrically on the left and right roof planes.

User: Does the house include a crawling space?

SpatialRGPT: Unknown, GT=Yes

There's no visible vent grid or elevation suggesting a ventilated plinth. Sub-floor access may be internal or obscured in this point cloud angle.

User: Are attics living spaces included?

SpatialRGPT: No, GT>No

No dormer extensions or roof windows suggestive of habitable attic conversion are visible. The upper roof volume appears sealed, likely used for insulation or storage only.

User: What's the window count on the front facade?

SpatialRGPT: 6 Windows, GT=7

3 on the ground floor, and 3 on the upper floor. No side windows detected from this viewpoint.

# CoT Distillation by Imitation via CE Minimization

$$\mathcal{L}_{\text{DCoT}} = -\frac{1}{T} \sum_{t=1}^T \log p_{\theta_S}(y_t^{(T)} \mid y_{<t}^{(T)}, x_{mm})$$

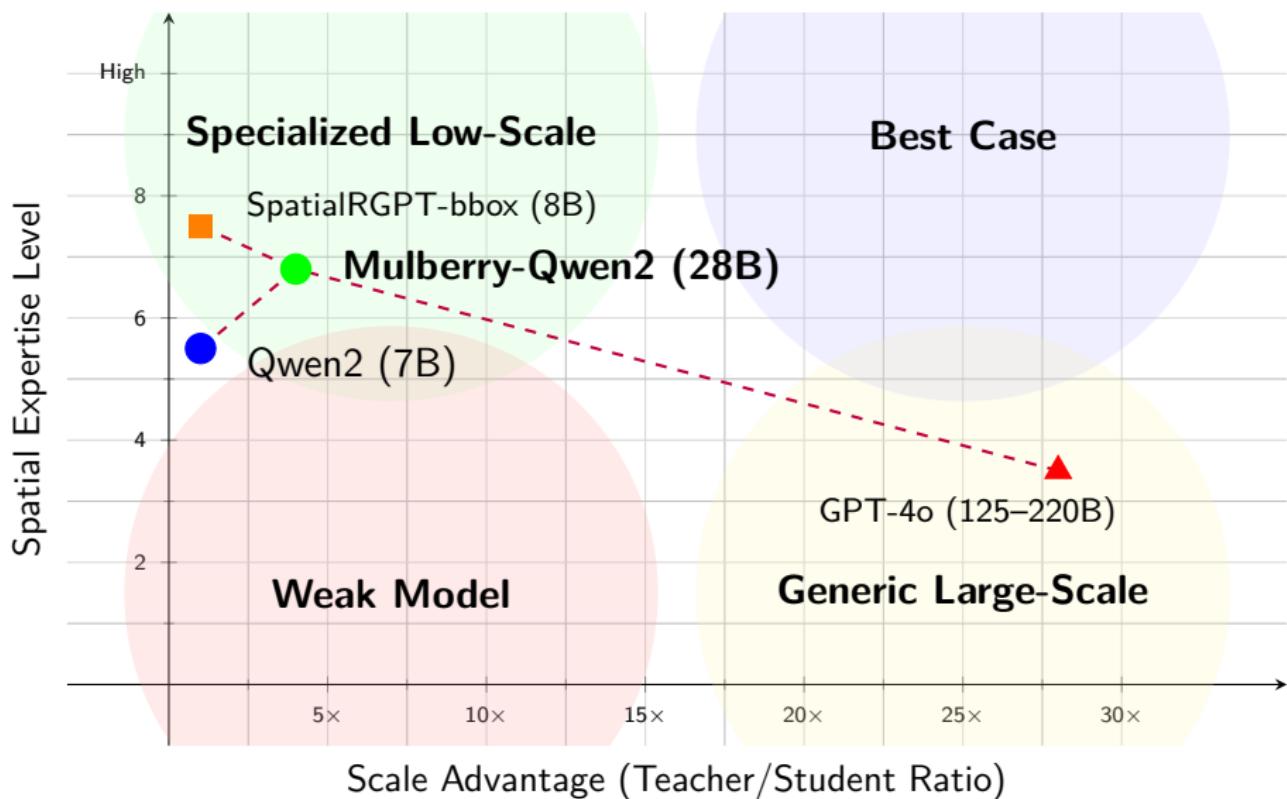
## Cross-Entropy vs Entropy Minimization

- **Assumption:** Correct Reasoning yields Low Entropy
- **Alignment:** Convergence & Overfitting
- **Specificity:** Penalizes incorrect but confident predictions
- **Correctness:** Abstract reasoning transfer

## Multimodal vs Distilled CoT

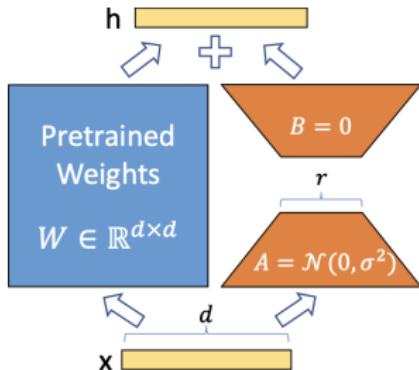
Configuration	MCoT	DCoT (Ours)
Train Data	Human-annotated rationales	Raw Teacher rationales
Train Size	29K (ScienceQA + A-OKVQA)	32 examples
Train Paradigm	SFT on CE Loss	KD via SFT on CE Loss
Train Components	Vision-language fusion layers	LLM attention modules only
Inference Process	Two-stage (rationale → answer)	Single-stage (end-to-end)
Parameter Efficiency	Full fine-tuning	LoRA
Reasoning Type	Multimodal	Complex chains

# Scaling Laws: Teacher-Student Knowledge Distillation

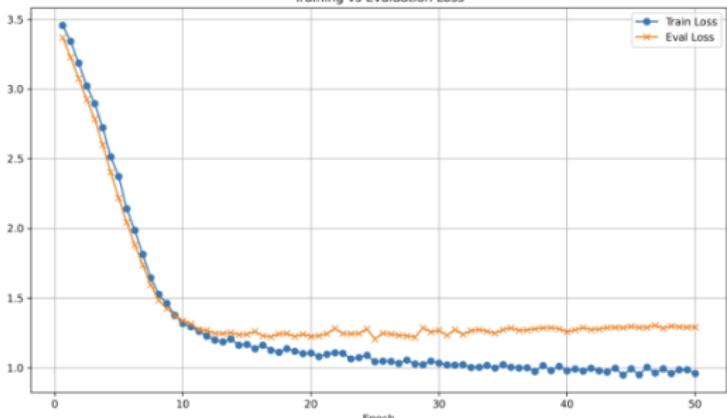
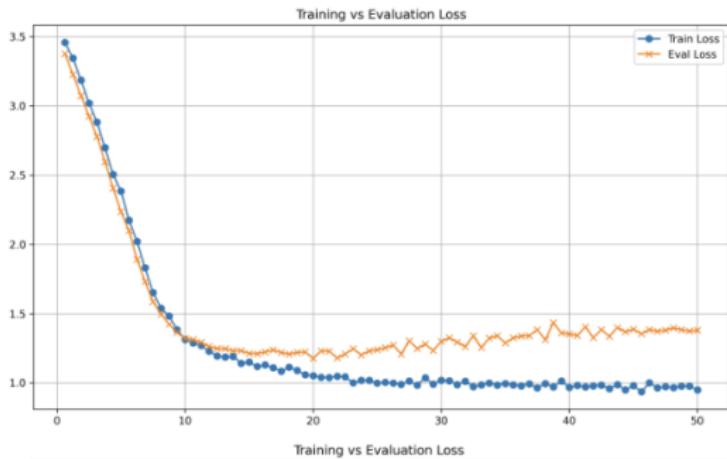


*Our methodology can be further scaled!*

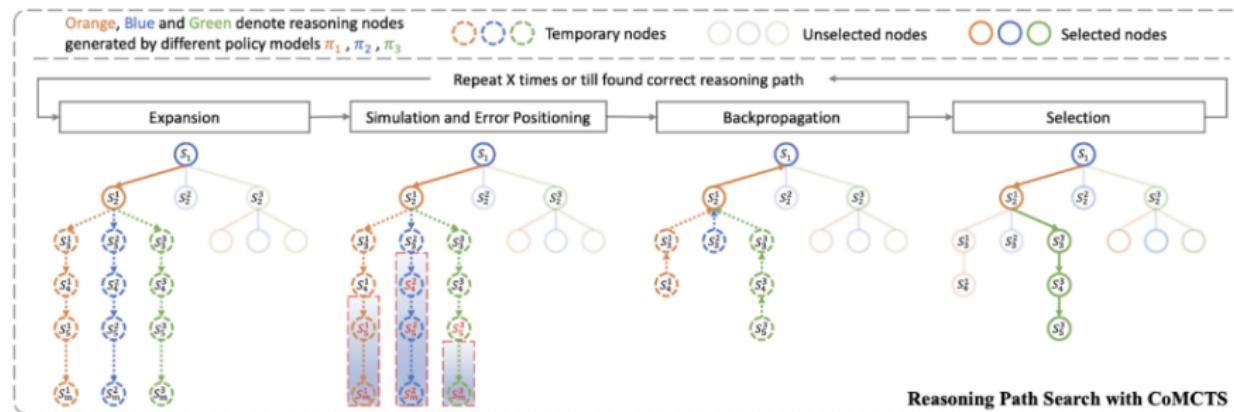
# Parameter-Efficient Fine-tuning with LoRA



- learning\_rate: 2e-5 (linear warmup)
- r (rank): 16
- lora\_alpha: 32
- early\_stopping: patience = 3
- target\_modules: ["q\_proj", "k\_proj", v\_proj", "o\_proj"]



# Test-Time Enhancement: Mulberry-Qwen CoMCTS



Configuration	Mulberry	Mulberry-Qwen (Ours)
Model Selection	GPT-4o, LLaMA, Qwen	Qwen-7B + Qwen-DCoT
Search Strategy	CoMCTS	CoMCTS
Training Strategy	SFT on Mulberry-260K	Zero-shot
Max Iterations	20	3
Reasoning Variation	High (inter-model)	Reduced (local/global)
Reasoning Robustness	Diverse patterns	Consistent domain-specific patterns
Bias Mitigation	Cross-model vote	Homogeneous model vote
Hallucination Reduction	High (multi-model)	Low (same-family)
Test Time Compute	Higher	Lower

# Qualitative Illustration of Qwen-CoMCTS Spatial Reasoning

Ground Truth: PV Panels Yes, Attic:No

**Question:** Does the house include photovoltaic panels and attics converted to living spaces?

Level 1

The house has a **pitched roof with visible elements**; user asks about rooftop and internal structural features.

Level 2

The roof contains **two distinct rectangular dark-colored units aligned symmetrically**- possibly solar panels. **No roof windows are visible.**

If the roof has no windows but has solar panels, **there may be no attic conversion**. Presence of PV suggests **energy infrastructure is present in attic area.**

Step 1: Count **2 solar panels**. Each - 1m x 1.7m. **Pitched on either side.**

Step 2: **No dormers, skylights, or visible attic access points** on the Roof.

Step 1: Presence of PV units → electric conduits likely routed to attic.

Step 2: Absence of dormers/windows → attic used for infrastructure

Level N

The final answer is PV = Yes, Attic = No

PV = Yes,  
Attic = No

PV = Unknown,  
Attic = Unknown

PV = Unknown,  
Attic = No

PV = Yes,  
Attic = No

# Evaluation Protocol & Metrics

## Experimental Design

- **Zero-shot:** All models on 45 images
- **Fine-tuned:** 4 Qwen variants on test set
- **Multiple runs:** 10 evaluations per configuration
- **Robustness:** Average  $\pm$  standard deviation

**Table:** Performance Metrics and Their Evaluation Focus

Metric(s)	Evaluation Focus
Accuracy, Balanced Accuracy	→ Generalization
Precision, Recall, F1-score	→ Hallucination Control
MAE / MSE	→ Counting Accuracy (e.g., windows)
Feature-wise Analysis	→ Spatial Reasoning via Feature Complexity

$$\text{F1-score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

# Results Overview: Model Performance

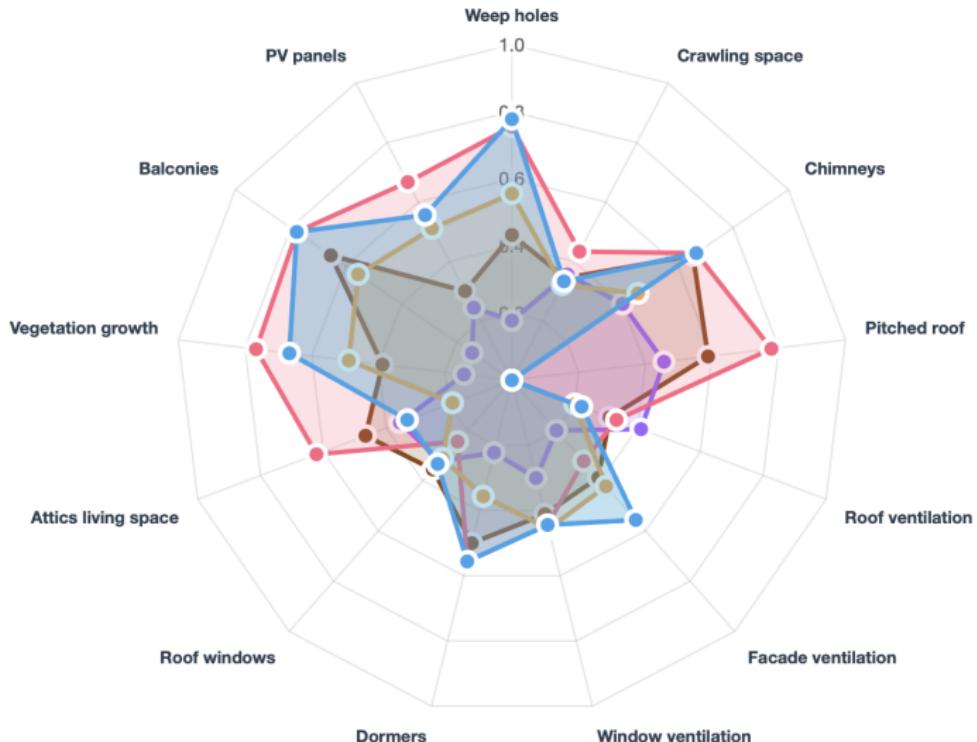
Table: Performance Comparison Across Models

Model	Accuracy	B. Acc.	Prec.	Recall	F1	MAE
GPT-4o*	<b>0.567</b>	<b>0.458</b>	<b>0.362</b>	<b>0.429</b>	<b>0.373</b>	<b>2.60 ± 0.14</b>
Mulberry-Qwen	0.479	0.370	0.178	0.333	0.226	3.78
Qwen2-VL-7B*	0.371	0.341	0.181	0.274	0.200	4.44
SpatialRGPT-bbox*	0.446	0.371	0.197	0.372	0.240	3.31 ± 0.29
SpatialRGPT*	0.287	0.330	0.139	0.311	0.171	6.67

## Key Findings

- **GPT-4o dominates** across all metrics
- **Mulberry-Qwen achieves 11% accuracy improvement** over baseline
- **Bounding box guidance helps SpatialRGPT significantly**
- **Gap with GPT-4o reduced:** 20%-8% in Acc., 11-8% in B. Acc.
- **Hallucinations reduced significantly** with DuTCh SpaCE
- **Robustness:** All models with (\*) showed std. dev. so Mulberry-Qwen is the most robust overall and feature-wise

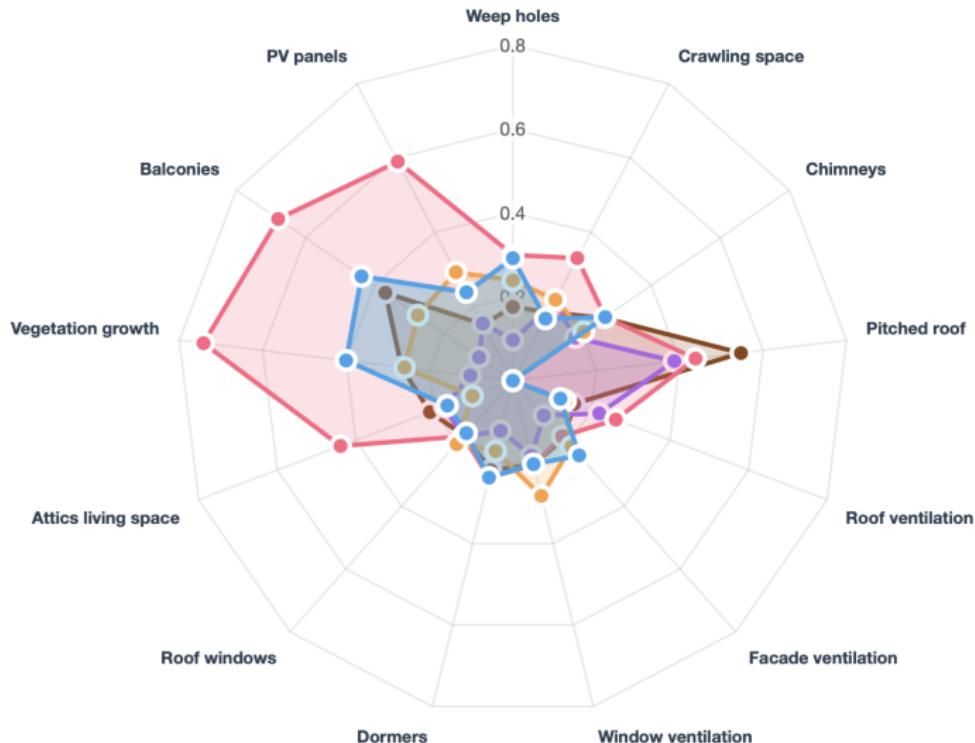
# Feature Performance (Accuracy)



● Mulberry-Qwen DCoT   ● Mulberry-Qwen DCoT+Scene   ● Qwen2-VL-7B Base   ● GPT-4o Base

● SpatialRGPT Base   ● SpatialRGPT Bbox

# Feature Hallucinations (F1-score)



● Mulberry-Qwen DCoT   ● Mulberry-Qwen DCoT+Scene   ● Qwen2-VL-7B Base   ● GPT-4o Base

● SpatialRGPT Base   ● SpatialRGPT Bbox

# Summary: DuTCh Space Achievements

MLLM Benefits Before	with DuTCh Space	Limitations After
<ul style="list-style-type: none"><li>Zero-shot transfer capabilities</li><li>Contextual reasoning about features</li><li>Interpretable explanations via CoT</li><li>Scalable to large building stocks</li></ul>	<ul style="list-style-type: none"><li>Knowledge distillation works</li><li>Test-time compute helps</li><li>Small models can compete</li><li>Net effect: Reduced under-prediction (FN), slight over-confidence (FP)</li></ul>	<ul style="list-style-type: none"><li>Spatial complexity challenges</li><li>Numerical reasoning gaps</li><li>Domain-specific biases</li><li>Computational requirements</li></ul>

## Key Takeaway

- Successful Mitigation:** DCoT + MC Tree Search reduces hallucinations by promoting evidence-based reasoning
- Reasoning compensates** for limited visual grounding
- Domain expertise** can rival raw scale

# RQ1 Findings: Model Architecture Comparison

## RQ1a: SoTA Comparison

**GPT-4o superior across all metrics** (20-30% gap with smaller models)

- Benefits from scale ( $\sim 125\text{-}220B$  vs. 7-8B parameters)
- Comprehensive multimodal training
- Better generalization to Dutch-specific features

## RQ1b: Bounding Box Guidance

**SpatialRGPT with bounding boxes:** 15% accuracy improvement

- Helps with spatial localization
- Reduces false positives
- Doesn't improve balanced accuracy significantly
- Additional annotation step

**Architecture Insights:** Different MLLMs have complementary strengths for different feature types

## RQ2 Findings: Enhancement Methods

### RQ2a: Scene Graph Augmentation

**No performance difference** between DCoT with/without scene graphs

- Qwen architecture may not effectively utilize explicit spatial representations
- DCoT reasoning already captures sufficient spatial relationships
- Need deeper integration beyond simple text augmentation

### RQ2b: LoRA Fine-tuning

**Significant improvements** across all metrics

- 10% accuracy improvement over baseline
- Reduces gap with GPT-4o from 20% to 8%
- Quality reasoning in few examples can match mitigate lack of extensive
- Maintains general capabilities while adding domain expertise

# Limitations & Challenges

## Dataset Limitations

- **Scale:** Only 45 images (vs. typical 100K+ datasets)
- **Annotation:** Subjective "unknown" categories
- **Imbalance:** Most features absent in dataset

## Technical Limitations

- **Visual Grounding:** No vision encoder fine-tuning
- **Spatial Integration:** Scene graphs not effectively utilized
- **Problem Complexity:** 14 questions in single prompt
- **Computational:** Limited test-time search iterations

## Methodological Considerations

- Cross-entropy loss vs. preference learning
- LoRA vs. full fine-tuning trade-offs
- Teacher quality vs. scale in distillation

# Future Work & Research Directions

## Immediate Extensions

- **Scale dataset:** 1000+ images from web scraping + automated filtering
- **Problem decomposition:** Split into building-section-specific questions
- **Vision Grounding:** Fine-tune multimodal components
- **Advanced search:** Increase CoMCTS iterations and model diversity

## Advanced Methodologies

- **RLHF:** Human preference optimization for spatial reasoning
- **Multi-adapter:** Feature-specific LoRA modules
- **Grounded CoT:** Visual evidence linking in reasoning chains
- **DoRA:** Weight decomposition for stable fine-tuning

# Broader Impact & Applications

## Energy & Sustainability

- Accelerate building renovation assessment
- Reduce manual inspection costs (time & labor)
- Enable large-scale retrofit planning
- Support EU Green Deal objectives

## Technical Contributions

- **Domain Adaptation:** Framework for specialized applications
- **Multimodal Reasoning:** Insights into reasoning vs. grounding
- **Knowledge Distillation:** Teacher-student dynamics in few-shot settings
- **Test-Time Compute:** Practical application in complex reasoning

# Conclusions

## DuTCh SpaCE key findings

- ① MLLMs are beneficial for facade analysis with proper enhancement
- ② Reasoning can compensate for limited visual grounding capabilities
- ③ Domain expertise rivals scale in knowledge transfer scenarios

## Theoretical Contributions

- First systematic MLLM evaluation on few-shot Dutch architecture
- Novel dual-teacher Distillation Framework for Reasoning vs Grounding
- Scaling laws in few-shot multimodal domain adaptation

## Practical Impact

- Accessible path to spatial reasoning enhancement
- Framework for domain-specific MLLM adaptation
- Cost-effective alternative to explicit spatial grounding
- Foundation for automated building assessment systems

# Acknowledgements

## Supervision & Guidance

- **Dr. Metehan Doyran** - First Examiner, Utrecht University
- **Dr. Itir Önal Ertugrul** - Second Examiner, Utrecht University
- **Raphaël Gueulet** - External Supervisor, TNO Machine Learning Engineer

## TNO Research Support

- **Paolo de Heer** - Data Scientist, **Wietske van Kanten-Roos & J.M. Tang** - Research managers

## Personal Support

- **Family** - Unwavering support throughout studies
- **Friends** Aryan Ashar, Shang-Jen Wang, Tom Slik - IT Support
- **Sarthak Anand** - Chat template implementation advice

*Thank you for making this research possible!*

# Thank You!

---

## Questions & Discussion

---

**Riccardo Campanella**

r.campanella@students.uu.nl

Utrecht University, MSc Artificial Intelligence

*Spatial Reasoning in Multimodal LLMs  
via CoT Distillation and Monte Carlo Tree Search  
for Dutch Facade-Element Detection: An Exploratory Study*